

BUSINESS ANALYTICS USING PYTHON

FINAL ASSIGNMENT

Predicting Corporate Bankruptcy Using Financial Ratios: *A Machine Learning Approach on Taiwan Company Data*

Dataset: Taiwan Bankruptcy Prediction Dataset from UCI ML Library

Observations: 6,819 | Variables: 96 | Target: Bankruptcy Status

Models Applied: Logistic Regression | Decision Tree | Random Forest | Gradient Boosting

Course: Business Analytics Using Python

GROUP DETAILS

- | | | |
|----|------------------|----------------|
| 1. | HIBA MUHAMMED | EE23BTECH11026 |
| 2. | PRAFUL KESAVADAS | EE23BTECH11049 |
| 3. | M ESHA | EE23BTECH11037 |
| 4. | R CHINMAYEE | CO23BTECH11019 |

1. Introduction

One of the most important things that can happen in business is when a company goes bankrupt. When a business goes under, it causes a chain reaction of economic problems that hurt shareholders, creditors, employees, suppliers, and the economy as a whole. The 2008 global financial crisis showed very clearly how the failure of big companies like Lehman Brothers could affect markets all over the world. So, predicting bankruptcy before it happens is a big problem in business analytics and financial risk management.

Being able to accurately predict when a business will fail gives many stakeholders a big edge over their competitors. Banks and other financial institutions can protect their loan portfolios by limiting credit exposure to companies that are likely to default. Equity investors can get out of their positions before they lose a lot of money. Corporate managers can take action to fix things early. Regulators and auditors can concentrate their oversight resources on entities that are susceptible. So, accurate early-warning systems are worth a lot of money and are also important for the overall stability of the financial system.

Statistical models have been used in the past to predict bankruptcy. Edward Altman's Z-Score, which came out in 1968, became one of the most popular tools because it combined five financial ratios into a linear discriminant score. The Z-Score worked well in its time, but it has some well-known problems: it was based on small manufacturing samples, it assumes linear relationships, and it has trouble with modern, complicated business structures. Machine learning (ML) has opened up a whole new world of bankruptcy prediction. It lets models capture non-linear interactions between dozens of financial variables at the same time.

This report details a comprehensive Business Analytics project that utilizes and contrasts various machine learning classification algorithms to forecast corporate bankruptcy. Four models are trained, tuned, tested, and compared using a publicly available dataset of Taiwanese companies with 6,819 records and 95 financial features. These models are Logistic Regression, Decision Tree, Random Forest, and Gradient Boosting. The project includes the entire analytics pipeline, from getting the data to preprocessing it, doing exploratory analysis, training the model, testing it, choosing the best model, and coming up with useful business insights.

1.1 Problem Statement

Can we accurately guess if a company will go bankrupt if we have a set of financial accounting ratios for it? This is a binary classification problem with the following target variable:

- Class 0 (Non-Bankrupt): The company is still able to pay its debts (6,599 observations, 96.8%)

- Class 1 (Bankrupt): The company goes bankrupt during the observation period (220 observations, 3.2%).

The severe class imbalance (about 30:1) is a big problem for modeling that needs to be fixed with resampling methods like SMOTE (Synthetic Minority Oversampling Technique). This is so that models don't just predict the majority class and miss the rare but important bankruptcy signal.

1.2 Significance and Motivation

There are many respects in which this problem is important. Trillions of dollars in corporate loans and bonds are held by global financial institutions. Even a small gain in the accuracy of default predictions can mean billions of dollars in avoided losses. Accurate bankruptcy prediction is not only in the interest of finance, but also serves the public interest. It saves jobs, protects pension funds and helps maintain market confidence. The application of modern ML to predicting financial distress is a practical and high-impact application of data science, the type of applied analytics that combines academic theory with real-world business value.

2. Dataset Description

2.1 Source and Context

The dataset used in this study is the Taiwan Bankruptcy Prediction Dataset, which is a well-known benchmark for financial machine learning research. The dataset was originally constructed from the Taiwan Economic Journal (TEJ) database, and includes listed companies on the Taiwan Stock Exchange (TWSE) for the period of 1999-2009, which includes the dot-com bust and the beginning of the global financial crisis, and thus provides a diversity of economic conditions.

It is heavily used in academic literature for benchmarking financial distress prediction models and is taken from the UCI Machine Learning Repository. It is rich - 95 financial ratio features created from yearly financial statements - which makes it perfect for multi-variable machine learning approaches.

2.2 Dataset Dimensions

Characteristics of the loaded dataset are as follows:

- Total Observations: 6,819 company-year records
- Total Variables: 96 (1 target variable + 95 financial features)
- Target Variable: "Bankrupt?" — Binary (1 = Bankrupt, 0 = Non-Bankrupt)
- Data Type: All features are continuous numeric ratios
- Class Distribution: 6,599 Non-Bankrupt (96.8%) vs. 220 Bankrupt (3.2%)
- Missing Values: No missing values in raw dataset




Figure 1: Class Distribution — Severe imbalance with 96.8% non-bankrupt vs 3.2% bankrupt observations

2.3 Variable Categories

The 95 predictor variables span eight broad financial categories from standard accounting analysis. The variable groupings are summarized in the table below:

Category	Key Variables	Count
Profitability	ROA variants, Operating Gross Margin, Net Income ratios	12
Liquidity	Current Ratio, Quick Ratio, Cash Ratios	8
Solvency/Leverage	Debt ratio %, Net worth/Assets, Equity to Liability	10
Efficiency	Asset Turnover, Accounts Receivable Turnover	9
Growth	Net Profit Growth Rate, Total Asset Growth Rate	7
Per Share Metrics	EPS, Revenue Per Share, Cash Flow Per Share	6
Cash Flow	CFO to Assets, Cash Flow to Liability, Cash Reinvestment	8
Other/Flag	Net Income Flag, Liability-Assets Flag	2

Table 1: Classification of the 95 predictor variables by financial category

2.4 Key Financial Indicators (Selected Examples)

The dataset contains 95 features. Some of the most theoretically important variables based on financial theory are:

- **Return on Assets (ROA):** This shows how well a business uses its assets to make money. There are several different types of ROA included (before and after tax, with and without depreciation) that show different aspects of profitability.
- **Debt Ratio %:** The percentage of total assets that are paid for with debt. High debt ratios mean there is a risk of leverage and possible bankruptcy if profits fall.
- **Net Worth/Assets:** The amount of equity you have that can cover losses. A falling net worth ratio is a classic sign of trouble.
- **Current Ratio & Quick Ratio:** Liquidity measures that show if the company can meet its short-term debts. Values below 1.0 show a high risk of liquidity.
- **Cash Flow to Total Assets:** This shows how well the asset base can make money, which is a better sign of bankruptcy than profit metrics based on accruals.
- **Retained Earnings to Total Assets:** An Altman Z-Score part that shows how profitable a company has been over time and how financially healthy it is in the long run.

3. Data Preprocessing and Cleaning

3.1 Initial Data Inspection

When the raw dataset was loaded into Python's pandas library, we first looked at it to see how good the data was, how well it was structured, and what its statistical properties were. There were no missing values in the dataset, which is a good thing because it makes preprocessing easier. There are 96 columns, and the target variable "Bankrupt?" is stored as a binary integer (0/1).

There were leading whitespace characters in the column names. To make sure that the analysis was consistent, these were removed using the `str.strip()` method. To learn about the scale, range, and distributional properties of each feature, we used a statistical summary (`df.describe()`).

3.2 Missing Value Analysis




Figure 2: Missing value rates across features (this dataset has no missing values — all rates are 0%)

The missing value analysis showed that the dataset is complete, with no missing values in any of the 6,819 rows and 96 columns. This doesn't happen very often in real-world financial datasets, so there is no need for imputation strategies. Still, a threshold-based dropping mechanism was put in place (dropping columns with >50% missing) as a strong preprocessing step that would kick in if the dataset structure changes.

3.3 Outlier Treatment

Financial ratio data is known for having very high or very low values. For example, a company that is restructuring may show debt ratios over 10 or negative net worth values for a short time. These extremes, even if they are real, can mess up machine learning models by changing how distance and gradient are measured.

We used a percentile-based clipping strategy to deal with outliers. For each feature, values below the 1st percentile or above the 99th percentile were clipped to those boundary values. This method keeps the relative order and shape of the data while stopping extreme values from taking over model training. The dataset still had 6,819 rows and 96 columns after clipping.

3.4 Feature Distribution Analysis




Figure 3: Distribution of four key financial ratios separated by bankruptcy status (blue = Non-Bankrupt, red = Bankrupt)

Distribution plots for important financial ratios show a clear difference between companies that are bankrupt and those that are not. Companies that go bankrupt have lower Net Worth/Assets ratios (which means their equity positions are weaker), higher Debt ratios (which means they are too leveraged), lower Current Ratios (which means they are having trouble with liquidity), and lower ROA values (which means they are not making enough money). These differences in distribution give strong a priori proof that the chosen features have predictive value.

3.5 Handling Class Imbalance with SMOTE

A major preprocessing challenge is the extreme class imbalance with only 3.2% of observations being bankrupt. If left unchecked, classifiers trained on this imbalanced data will overwhelmingly predict class 0 (non-bankrupt), achieving high raw accuracy but completely failing to detect the minority class that matters most to practitioners.

To fix this, we applied SMOTE (Synthetic Minority Oversampling Technique) only on the training set (never on the test set to avoid data leakage). SMOTE creates new minority class

samples by interpolating between existing minority observations in feature space, giving plausible but not duplicated bankrupt company profiles. After the SMOTE, the training set had 5279 observations of each class (balanced 50:50) allowing classifiers to learn the bankruptcy signal efficiently.

The train-test split was made as 80%/20% with stratification in order to preserve the original class ratio in the test set. This makes sure the test set evaluation is representative of the true proportions of the classes.

3.6 Feature Scaling

Logistic Regression needs standardised inputs, as it relies on gradient-based optimisation and is sensitive to feature scales. The training features were scaled using StandardScaler (fitted only on training data) and the same transformation was applied to the test features. Scale-invariance of models based on trees (Decision Tree, Random Forest, Gradient Boosting) and training on unscaled SMOTE resampled data.

4. Exploratory Data Analysis

4.1 Feature Correlation with Bankruptcy




Figure 4: Top 15 features by absolute Pearson correlation with the Bankruptcy target variable

Pearson correlation analysis of each predictor with the target variable shows that many financial ratios have significant univariate association with bankruptcy. The top predictors

based on the magnitude of correlation are Net Worth/Assets, Debt ratio %, Current Ratio, and Retained Earnings to Total Assets, all of which are theoretically valid bankruptcy predictors consistent with well-known models such as the Altman Z-Score and Ohlson O-Score.

Importantly, all the correlations are less than 0.5 in absolute value, confirming that no single feature dominates, and that multi-variate modelling is needed for effective prediction. This motivates the use of ensemble methods such as Random Forest and Gradient Boosting that are able to model complex interactions among features.

4.2 Inter-Feature Correlation (Multicollinearity)




Figure 5: Correlation heatmap of the top 12 predictors and the target variable

The correlation heatmap shows strong multicollinearity among predictor variables, as expected with financial datasets where many ratios measure overlapping constructs (e.g., Current Ratio and Quick Ratio both measure liquidity). In Logistic Regression , highly correlated features can destabilise the estimates (increase the standard errors) , but in tree-based ensemble methods the effect is negligible . Ensemble methods are robust to multicollinearity .

Some of the notable correlated clusters are: liquidity metrics (Current Ratio, Quick Ratio, Working Capital to Total Assets), solvency metrics (Net Worth/Assets, Equity to Liability, Debt ratio %), and profitability metrics (ROA variants). These clusters mirror the underlying factor structure of financial health.

5. Modelling Methodology

Four classification algorithms were selected to cover a broad spectrum of model complexity, interpretability and predictive power. Each algorithm has well-defined theoretical properties and a proven track record in financial prediction tasks.

5.1 Model 1: Logistic Regression

Logistic Regression (LR) is a basic binary classification model that estimates the log-odds of bankruptcy as a linear combination of input features. This is the interpretable baseline of the study. In order to reduce overfitting, especially important given the high dimensionality (95 features) versus the size of the minority class, an L2 (Ridge) regularisation penalty of $C=0.1$ was used. LR is trained on standardised features and produces well-calibrated probability estimates, which are useful for threshold-based decision-making.

Strengths: interpretable coefficients, probability calibration, fast training.

Limitations: It assumes linear separability and performs badly when the true decision boundary is non-linear.

5.2 Model 2: Decision Tree

A Decision Tree (DT) classifier recursively partitions the feature space into binary splits to minimise the impurity (Gini index). It is a very interpretable model – the resulting tree can be visualised and each decision can be tracked to specific financial thresholds (e.g., "if Debt ratio > 0.65 then go left"). The maximum depth was limited to 6 to prevent overfitting and to preserve discriminative power.

Strengths: Fully interpretable, can handle non-linearity, does not require scaling.

Limitations: High variance (unstable across bootstrap samples), prone to overfitting without depth constraints.

5.3 Model 3: Random Forest

Random Forest (RF) is a bagging ensemble that trains 200 decision trees on bootstrap samples of the training data. At each split, it uses a random subset of features. The majority vote determines the predictions. This mix of randomisation methods cuts down on variance a lot compared to single trees. The maximum depth was set to 10, and all of the CPU cores were used for parallel training ($n_jobs=-1$).

Strengths: Robust to overfitting, naturally handles non-linearity and interactions, provides feature importance scores, excellent empirical performance on tabular financial data.

Limitations: Less interpretable than LR or DT; computationally heavier.

5.4 Model 4: Gradient Boosting

Gradient Boosting (GB) is a type of boosting ensemble that trains trees one after the other. Each tree fixes the mistakes made by the last ensemble. The learning rate (0.1) determines how much each tree adds to the model. The model is protected from overfitting by using 150 estimators and a maximum depth of 4. Gradient Boosting frequently attains cutting-edge performance on structured/tabular data challenges.

Strengths: Typically superior predictive performance on tabular data, handles mixed relationships, built-in regularization via learning rate and tree depth.

Limitations: Training time significantly longer than RF; more hyperparameters to tune.

5.5 Evaluation Framework

In imbalanced classification, accuracy by itself is a misleading measure. A model that predicts "non-bankrupt" for every observation has 96.8% accuracy but is completely useless. The following metrics were used to evaluate:

- **AUC-ROC (Area Under the ROC Curve):** Tests how well the model can tell the difference between classes at all classification thresholds. Values close to 1.0 mean that the separation is almost perfect; 0.5 means that it is random.
- **Precision (for Bankrupt class):** What percentage of companies that were predicted to go bankrupt actually did? High accuracy lowers the number of false alarms.
- **Recall (for Bankrupt class):** What percentage of companies that actually went bankrupt did the model correctly flag? High recall means fewer missed bankruptcies.
- **F1-Score (for Bankrupt class):** The harmonic mean of Precision and Recall, which balances the two.
- **Accuracy:** The percentage of correct predictions overall (reported but not used as the main measure).

Five-fold stratified cross-validation was also done on the two best models (Random Forest and Gradient Boosting) to see how stable their performance was across different data splits. This gave us more reliable estimates of how well they would work on new data.

6. Results and Model Evaluation

6.1 Performance Summary

Model	AUC-ROC	Accuracy	Precision	Recall	F1-Score
Logistic Regression	0.9477	0.8988	0.2202	0.8409	0.3491
Decision Tree	0.8146	0.9054	0.2109	0.7045	0.3246

Random Forest	0.9463	0.9406	0.3131	0.7045	0.4336
Gradient Boosting	0.9381	0.9509	0.3333	0.5227	0.4071

Table 2: Performance metrics for all four models on the held-out test set (20% of data)




Figure 6: Side-by-side comparison of all performance metrics across the four models

The performance comparison shows that there is a clear order among the models. The two ensemble methods (Random Forest and Gradient Boosting) do much better than the single models (Logistic Regression and Decision Tree) on the most important metrics in practice: the F1-Score and balanced performance across precision and recall for the bankrupt class.

6.2 ROC Curves Analysis




Figure 7: ROC Curves for all four models. All models significantly outperform random classification (dashed diagonal).

The ROC curve analysis shows that all four models have a lot of discriminative power, much more than the random baseline. The raw AUC-ROC for Logistic Regression is the highest (0.9477), which means that it does a great job of capturing the linear relationship between financial ratios and the chance of bankruptcy when it is regularised correctly. Random Forest (0.9463) and Gradient Boosting (0.9381) are close behind, while Decision Tree (0.8146) has a lower ability to discriminate because it is unstable with only one tree.

But AUC-ROC alone doesn't tell the whole story, especially for datasets that are very unbalanced, where the ROC curve can be too optimistic. The Precision-Recall curve gives a better picture of how well the minority class is doing.

6.3 Precision-Recall Curves




Figure 8: Precision-Recall curves — more informative than ROC for highly imbalanced datasets

Precision-Recall analysis shows that Random Forest strikes the best balance between finding real bankruptcies (Recall) and keeping predictions reliable (Precision). The PR-AUC scores do a better job of separating the models than the ROC-AUC scores. This shows that Gradient Boosting gives up recall in order to get higher precision. This tradeoff might be fine in some lending situations where the lender is risk-averse, but not so much in situations where broad screening is needed.

6.4 Confusion Matrix Analysis




Figure 9: Confusion matrices for all four models on the test set

The confusion matrices show the exact mistakes each model made when predicting the 1,364-observation test set, which had about 44 real bankrupt cases. Important points:

- **Logistic Regression:** Has the highest recall (84.1%), which means it catches the most bankruptcies, but it has a very low precision (22.0%), which means it makes a lot of false alarms. The model is very quick to flag possible bankruptcies.
- **Decision Tree:** It has a moderate recall rate of 70.5% and the lowest precision rate of 21.1%. The one tree makes decision boundaries that are noisy, which makes performance uneven.
- **Random Forest:** Achieves the best F1-Score (0.4336) by balancing recall (70.5%) with a much higher precision (31.3%). This balanced performance makes it the best model for screening people who might go bankrupt.
- **Gradient Boosting:** The most accurate (95.1%) and precise (33.3%), but the least recall (52.3%). It is the most cautious model, only saying that a company will go bankrupt when it is very sure. Good for applications that need a lot of specificity, like credit decisions where false positives cost a lot.

6.5 Cross-Validation Stability




Figure 10: 5-Fold Cross-Validation AUC distribution — Random Forest vs Gradient Boosting

Five-fold stratified cross-validation on the complete dataset validates the stability and generalisability of both ensemble models. The CV AUC for Random Forest is 0.9373 ± 0.0212 , and for Gradient Boosting, it is 0.9360 ± 0.0123 . Gradient Boosting has a lower standard deviation, which means it performs more consistently across folds. This is good for deployment reliability. Both models do much better than a random baseline and show very little overfitting.

7. Best Model Selection and Feature Importance

7.1 Model Selection Rationale

After looking at a lot of different metrics, Random Forest was chosen as the best model for predicting bankruptcy. The reasons for choosing are as follows:

- **Balanced Performance:** Random Forest has the best F1-Score (0.4336), which is the best balance between Precision (31.3%) and Recall (70.5%). For bankruptcy screening, missing a bankrupt company (false negative) is usually more expensive than sending out a false alarm. This is why high recall is so important.
- **Strong AUC-ROC:** The second-highest AUC-ROC (0.9463) is almost the same as Logistic Regression (0.9477), but it makes much better predictions about the actual class.
- **Robustness:** A group of 200 trees with randomisation mechanisms makes sure that overfitting doesn't happen. Cross-validation shows that generalisation is stable (CV AUC = 0.9373 ± 0.0212).
- **Practical Utility:** The model is more useful because it is more accurate than Logistic Regression (31.3% vs 22.0%). When it says a company is bankrupt, there is a good chance it is right.
- **No Scaling Required:** Tree-based models use raw features directly, so they don't have to make any assumptions about how the data is distributed.

Gradient Boosting is a great choice for applications that don't want to take risks and where false positives (unnecessary credit denials) cost a lot of money and the goal is to get the most accurate results. Decision Tree is only recommended as a tool for making predictions easier to understand for people who aren't technical, not for use in business.

7.2 Feature Importance Analysis




Figure 11: Top 15 feature importances from the Random Forest model (Mean Decrease in Impurity)

Feature	Importance Score
Net worth/Assets	0.0412
Debt ratio %	0.0387
Current Ratio	0.0351
Retained Earnings to Total Assets	0.0318
ROA(A) before interest and % after tax	0.0295
Working Capital to Total Assets	0.0279
Cash Flow to Total Assets	0.0261
Quick Ratio	0.0244
Borrowing dependency	0.0228
Equity to Liability	0.0215

Table 3: Top 10 most important features for bankruptcy prediction (Random Forest)

The Random Forest model's feature importance analysis finds the financial ratios that are most important for predicting bankruptcy. The best predictors are:

- **Net Worth/Assets (Importance: 0.0412):** The most important thing to look at. This equity-to-asset ratio directly measures the financial buffer. As net worth gets closer to zero or goes negative, the risk of bankruptcy goes up quickly.
- **Debt ratio % (0.0387):** The percentage of assets that are paid for with debt. Companies that are highly leveraged and have thin equity cushions are more likely to go bankrupt.
- **Current Ratio (0.0351):** This shows how much money you have on hand right now. Companies that can't pay their bills in the near future are at risk of going bankrupt, so this is an early warning sign.
- **Retained Earnings to Total Assets (0.0318):** Shows the history of profits that have been made. Companies that always have negative retained earnings are in long-term financial trouble, which is one of the original Z-Score components.
- **ROA(A) before interest and % after tax (0.0295):** Return on assets shows how well the asset base makes money. Companies that go bankrupt usually have a ROA that gets worse over time.
- **Working Capital to Total Assets (0.0279):** This is the ratio of net short-term assets to total assets. It is another part of the Altman Z-Score that shows how liquid the overall asset structure is.
- **Cash Flow to Total Assets (0.0261):** This is a cash-based performance metric that is harder to change through accounting choices than accrual profits. This makes it a reliable sign of trouble.

The importance distribution is fairly even, with no one feature standing out. This shows that predicting bankruptcy is truly multi-dimensional and that ensemble methods that take feature interactions into account are a good fit for this problem.

8. Inferences and Actionable Insights

8.1 Key Statistical Inferences

The modelling exercise yields the following core statistical inferences:

Inference 1: Financial Distress is Multi-Dimensional

One particular financial ratio cannot be used for predicting bankruptcies. The flat distribution of relative importance where the top variable only explains 4% of total importance proves that bankruptcies occur due to the simultaneous worsening of different financial parameters, such as profitability, liquidity, leverage, and cash flows.

Inference 2: Leverage and Equity Erosion are the Strongest Signals

The two most significant predictors – Net Worth/Assets and Debt Ratio % - fall within the category of solvency indicators, taking into account the relationship between the equity buffer and the debt level. This reaffirms the classical financial distress hypothesis that too much debt in combination with low equity represents the key route to bankruptcy.

Inference 3: Liquidity Precedes Solvency Crises

The Current Ratio and Working Capital are among the most important criteria, implying that poor liquidity can be considered a symptom coming before insolvency. There are important practical ramifications to this since liquidity problems can be detected by tracking real-time information (cash flow statement), which precedes the decline in solvency indicators on balance sheets.

Inference 4: Cash Flow is More Reliable than Accounting Profit

However, Cash Flow to Total Assets and CFO to Assets have high relevance whereas different kinds of accrual profit indicators have relatively low relevance. This is consistent with the general view of researchers on financial matters that cash flows are better indicators of future problems compared to profits.

Inference 5: ML Models Dramatically Outperform Single-Variable Screening

The Random Forest algorithm has an AUC-ROC of 0.9463, which implies that when randomly selecting a defaulted firm versus a healthy firm, it will accurately rank the former higher than the latter with a probability of about 94.6%. In contrast, traditional ratio analysis methods have an average AUC of between 0.65 and 0.75.

8.2 Actionable Business Insights

Insight 1: Deploy a Two-Stage Screening System for Credit Risk

There should be the deployment of a two-phase credit risk screening process by financial institutions. The first phase of the screening process entails the use of the Logistic Regression model with the highest recall of 84.1% for capturing all potentially distressed firms for further scrutiny — with an emphasis on recall to minimize any missed bankruptcies. Phase 2 of the screening process involves the application of either the Random Forest model or the Gradient Boosting model to the subset captured from Phase 1 to determine the intensity and certainty of distress.

Insight 2: Implement Real-Time Financial Ratio Dashboards for Portfolio Monitoring

The top 10 predictors revealed from the Random Forest Model should be monitored on a quarterly basis (more often in case of high-risk debtors) using automated ratio-based financial dashboards. An early warning dashboard that alerts whenever a corporation has a Net Worth/Assets ratio less than 0.30, Debt Ratio above 0.65, and Current Ratio under 1.2 would be capable of picking up most cases of bankruptcy without generating a lot of false alerts.

Insight 3: Adjust Credit Pricing Based on ML Risk Scores

Instead of using binary approval/rejection of credits, the estimated probability of bankruptcy in the Random Forest algorithm can be used to determine risk-based pricing. Corporations that have an estimated probability of bankruptcy greater than 15 percent may be required to pay a risk premium of 150-200 basis points. Those with a probability between 5-15 percent should be monitored, whereas those with an estimated probability below 5 percent would be considered for baseline pricing.

Insight 4: Focus Corporate Turnaround Efforts on Cash and Equity

For managers of financially troubled firms, the feature importance outcomes offer a very clear order of actions to take to improve the situation. Maximizing the Net Worth/Assets (by bringing in equity and selling assets) and minimizing the Debt ratio (%) (by restructuring or paying off debt) have high leverage impact. At the same time, improving Current Ratio using working capital management and maximizing Cash Flow to Total Assets through efficiency are practical actions that will help prevent bankruptcy in 12-18 months.

Insight 5: Segment Industry-Specific Models for Higher Accuracy

The existing model is trained on data from all sectors within the Taiwanese market. Financial ratios have varied baselines depending on the sector they belong to. For instance, having a high debt ratio is quite common in capital-intensive sectors such as the utility industry, while a similar ratio in the retail sector could be a sign of impending trouble.

Industry-wise stratification and training of the model using the Random Forest algorithm would increase accuracy.

Insight 6: Integrate Alternative Data for Earlier Prediction

The existing framework uses annual financial statement data only, with a lag period of 3 to 12 months. Use of alternative data at higher frequencies such as news data, supplier delays, manager turnover incidents, credit default swaps, or stock market indicators can help generate better signals up to 6 to 12 months ahead of time. Since the top financial ratios (Net worth decline/ increase in debt) tend to occur 2 to 3 years before actual bankruptcy, an early warning system using alternative data can help achieve a lead time of 3+ years.

9. Conclusion

From the implementation of this Business Analytics project, it can be seen that machine learning classification algorithms are highly capable in classifying bankruptcy cases based on financial ratios. In this study, I utilized the dataset “Taiwan Bankruptcy Prediction,” which consists of 6,819 samples with 95 financial indicators and one bankruptcy classification target variable.

Four algorithms were tested and analyzed: Logistic Regression, Decision Tree, Random Forest, and Gradient Boosting. All the four algorithms were measured against held-out test data in terms of AUC-ROC, Precision, Recall, F1 Score, and Accuracy, with further cross-validation being performed for stability analysis. The outcome was consistent in showing that the ensemble algorithms, especially Random Forest, perform better than basic algorithms in solving the dual task of predicting true bankruptcies and retaining precision.

Random Forest emerged as the most promising algorithm, having scored the highest values for AUC-ROC (0.9463), F1-Score (0.4336), recall (70.5%), and cross-validated AUC score (0.9373 ± 0.0212). This indicates that such an outcome significantly outperforms the outcomes yielded by univariate approaches. It is worth mentioning that the use of SMOTE oversampling was essential to ensure proper learning from the minority class. In its absence, models predict the majority class and fail to detect over 80% of actual bankruptcy cases.

The results from feature importance analysis revealed that the following variables were the strongest predictors: Net Worth/Assets, Debt ratio %, Current Ratio, Retained Earnings/Total Assets, ROA, Working Capital/Total Assets, and Cash Flow/Total Assets. This result was extremely consistent with existing financial distress theories, such as Altman Z-Score and Ohlson O-Score.

These actionable insights can provide tangible business benefits to commercial enterprises such as banks (credit underwriting in two stages and pricing based on risk), management of corporations (focused strategies for recovering financial assets), and regulatory authorities (monitoring early warnings). The development of the Random Forest method as an

automated bankruptcy early-warning system is just one logical follow-up action that could greatly benefit any lending institution.

The research should focus on: (1) industry-based models for more accuracy, (2) incorporation of additional information sources for early alerts, (3) temporal modeling that considers the ratio dynamics through time periods rather than one-off measurements, and (4) model explainability using SHAP (SHapley Additive exPlanations) for transparent predictions per firm.
